



**Carnegie Mellon**

**School of Computer Science**

# Efficient Parallel Learning of Linear Dynamical Systems on SMPs

*Lei Li*

Computer Science Department  
School of Computer Science  
Carnegie Mellon University  
leili@cs.cmu.edu

Motion stitching via effort minimization  
with James McCann, Nancy Pollard and  
Christos Faloutsos  
[Eurographics 2008]

Parallel learning of linear dynamical systems  
with Wenjie Fu, Fan Guo, Todd Mowry and  
Christos Faloutsos  
[KDD 2008]

# Background

- Motion Capture



- Markers on human body, optical cameras to capture the marker positions, and translated into body local coordinates.
- Application:
  - Movie/game/medical industry



# Outline

- Background
- Motivation: effortless motion stitching
- Parallel learning with Cut-And-Stitch
- Experiments and Results
- Conclusion

# Motivation

- Given two human motion sequences, how to stitch them together in a natural way( = looks natural in human's eyes)?



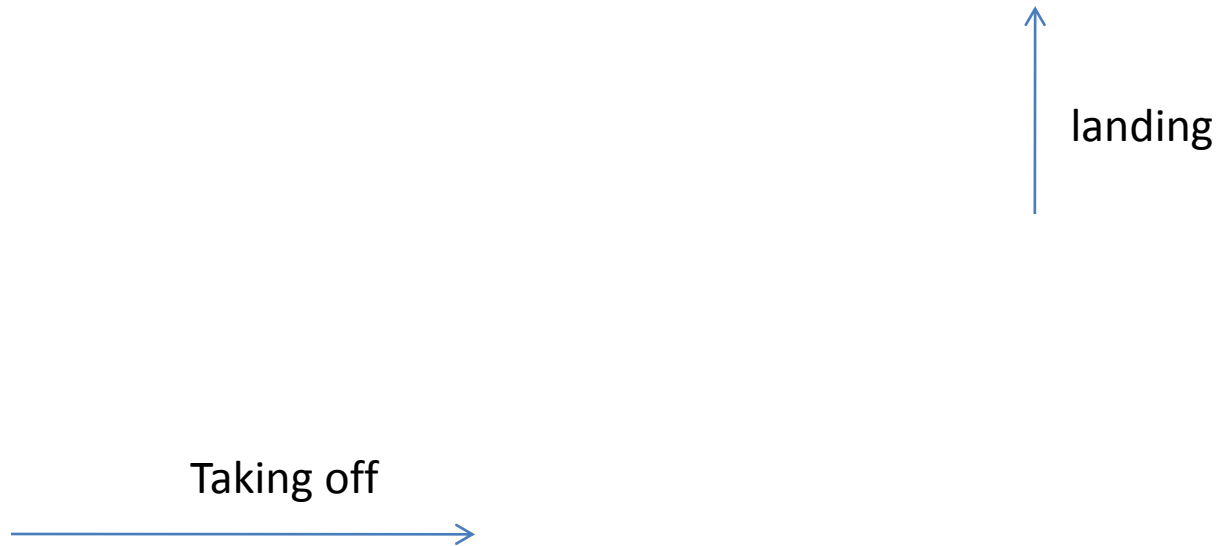
e.g. walking to running

- Given a human motion sequence, how to find the best natural stitchable motion in motion capture database?


# Intuition

- Intuition:
  - Laziness is a virtue. Natural motion use minimum energy
- Laziness-score (L-score) = energy used during stitching
- Objective:
  - Minimize laziness-score

# Example




# Example, Natural stitching





# But, how about this way?




# Observations

- Naturalness depends on smoothness
- Naturalness also depends on motion speed


# Proposed Method

- Estimate stitching path using Linear Dynamical Systems




# Proposed Method (cont')

- Estimate the velocity and acceleration during the stitching, compute energy (defined as L-score)





# Proposed Method (cont')

- Minimize L-score with respect to any stitching hops. (defined as elastic L-score)



# Example stitching



- [Link to video](#)



# Outline

- Background
- Motivation: effortless motion stitching
- Parallel learning with Cut-And-Stitch
- Experiments and Results
- Conclusion



# Parallel Learning for LDS

- Challenge:
  - Learning Linear Dynamical System is slow for long sequences
- Traditional Method:
  - *Maximum Likelihood* Estimation via *Expectation-Maximization*(EM) algorithm
- Objective:
  - **Parallelize** the learning algorithm
- Assumption:
  - *shared memory architecture*

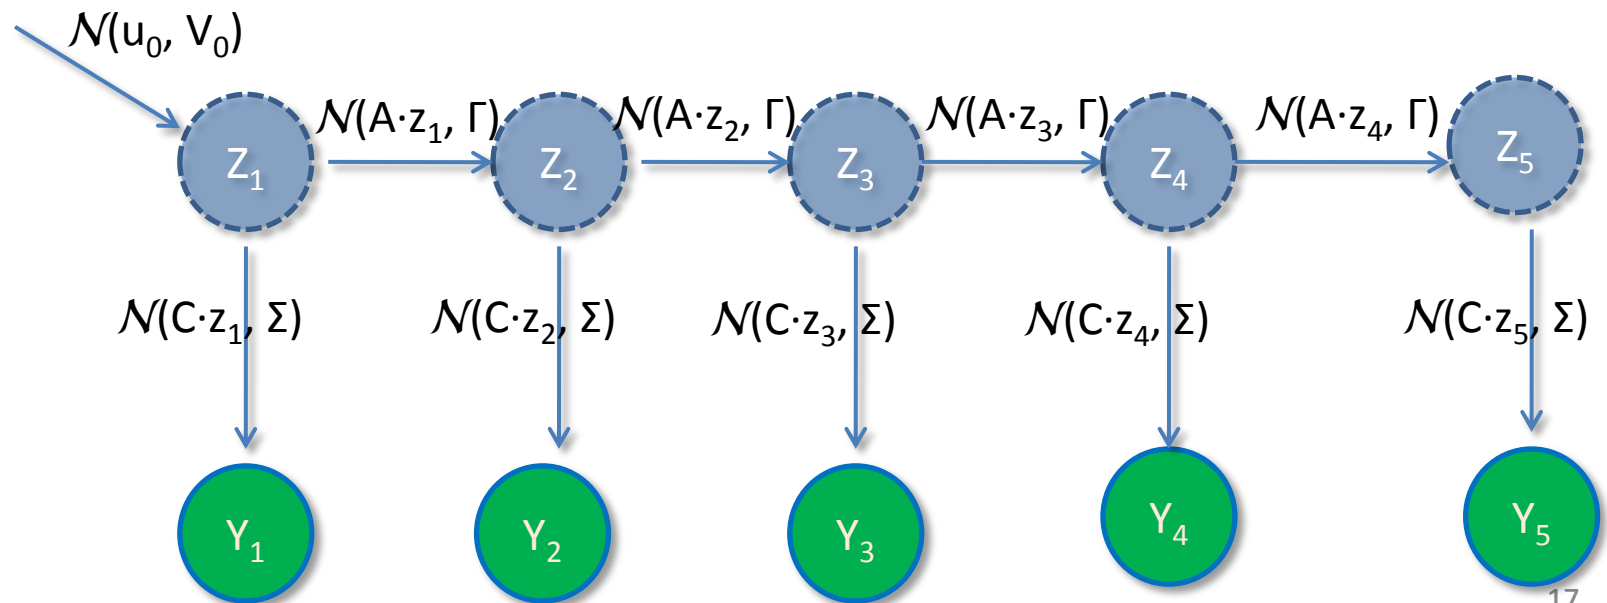




# Linear Dynamical System


*aka. Kalman Filter*

- Parameters:  $\theta = (u_0, V_0, A, \Gamma, C, \Sigma)$
- Observation:  $y_1 \dots y_n$
- Hidden variables:  $z_1 \dots z_n$






# Example



given positions, estimate dynamics (i.e. params)







# Traditional: How to learn LDS?



# Sequential Learning (EM)




Compute  $P(z_1 \mid y_1)$

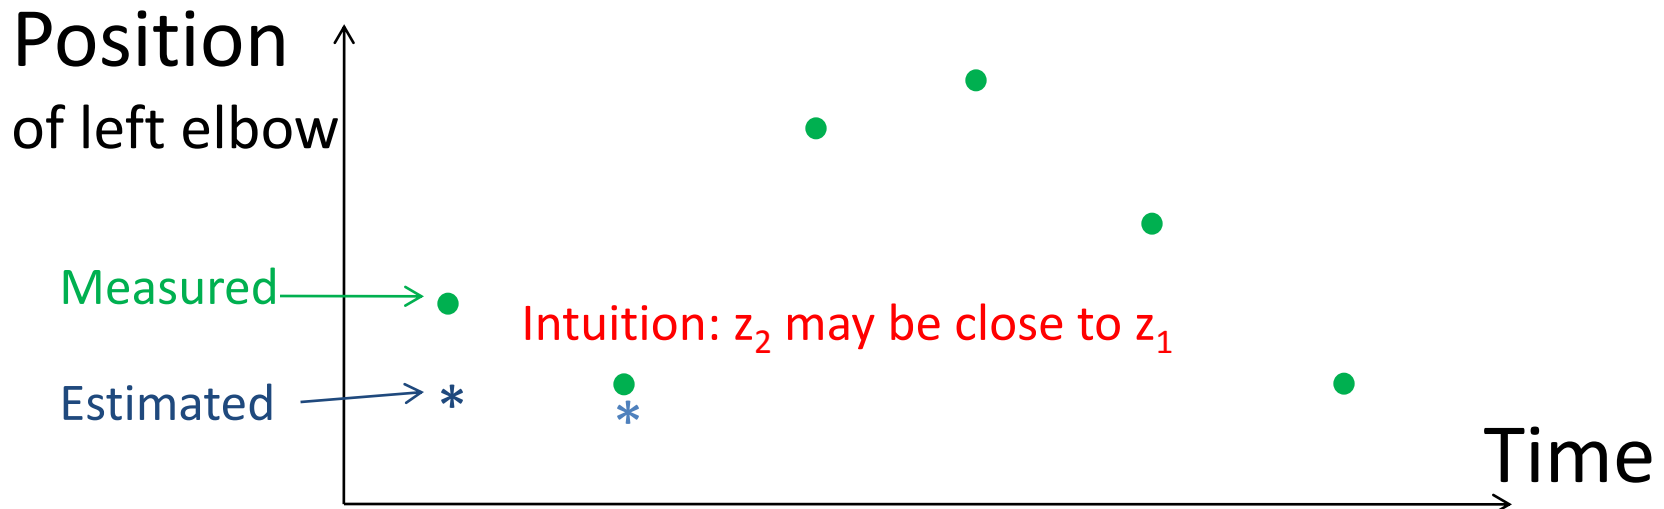




# Sequential Learning (EM)




From  $P(z_1 | y_1) \rightarrow$  Compute  $P(z_2 | y_1, y_2)$





# Sequential Learning (EM)




From  $P(z_2 | y_1, y_2) \rightarrow$  Compute  $P(z_3 | y_1, y_2, y_3)$

Position  
of left elbow


Measured  $\rightarrow$  ●

Estimated  $\rightarrow$  \*






# Sequential Learning (EM)




From  $P(z_3 | y_1, y_2, y_3) \rightarrow$  Compute  $P(z_4 | y_1, y_2, y_3, y_4)$






# Sequential Learning (EM)




From  $P(z_4 | y_1, y_2, y_3, y_4) \rightarrow$  Compute  $P(z_5 | y_1, y_2, y_3, y_4, y_5)$








# Sequential Learning (EM)




From  $P(z_5 | y_1, y_2, y_3, y_4, y_5) \rightarrow$  Compute  $P(z_6 | y_1, y_2, y_3, y_4, y_5, y_6)$






# Sequential Learning (EM)



From  $P(z_6 | y_1, y_2, y_3, y_4, y_5, y_6) \rightarrow$  Compute  $P(z_5 | y_1, y_2, y_3, y_4, y_5, y_6)$


Position  
of left elbow

Measured 






# Sequential Learning (EM)



From  $P(z_6 | y_1, y_2, y_3, y_4, y_5, y_6) \rightarrow$  Compute  $P(z_4 | y_1, y_2, y_3, y_4, y_5, y_6)$


Position  
of left elbow

Measured 






# Sequential Learning (EM)



From  $P(z_4 | y_1, y_2, y_3, y_4, y_5, y_6) \rightarrow$  Compute  $P(z_3 | y_1, y_2, y_3, y_4, y_5, y_6)$


Position  
of left elbow

Measured 






# Sequential Learning (EM)



From  $P(z_3 | y_1, y_2, y_3, y_4, y_5, y_6) \rightarrow$  Compute  $P(z_2 | y_1, y_2, y_3, y_4, y_5, y_6)$


Position  
of left elbow

Measured





# Sequential Learning (EM)




From  $P(z_2 | y_1, y_2, y_3, y_4, y_5, y_6) \rightarrow$  Compute  $P(z_1 | y_1, y_2, y_3, y_4, y_5, y_6)$

Position  
of left elbow

Measured


Estimated

Time





# Sequential Learning (EM)



From all posterior  $z_1, z_2, z_3, z_4, z_5, z_6$

$P(z_1 | y_1, y_2, y_3, y_4, y_5, y_6), P(z_2 | y_1, y_2, y_3, y_4, y_5, y_6) \dots$

Compute sufficient statistics


$$E[z_i]$$

$$E[z_i z_i']$$


$$E[z_{i-1} z_i']$$




# Sequential Learning (EM)



with sufficient statistics, compute  $\underset{\theta}{\operatorname{argmax}} \leftarrow \text{likelihood}(\theta)$







Speed Bottleneck:  
sequential computation of posterior

How to *parallelize* it?

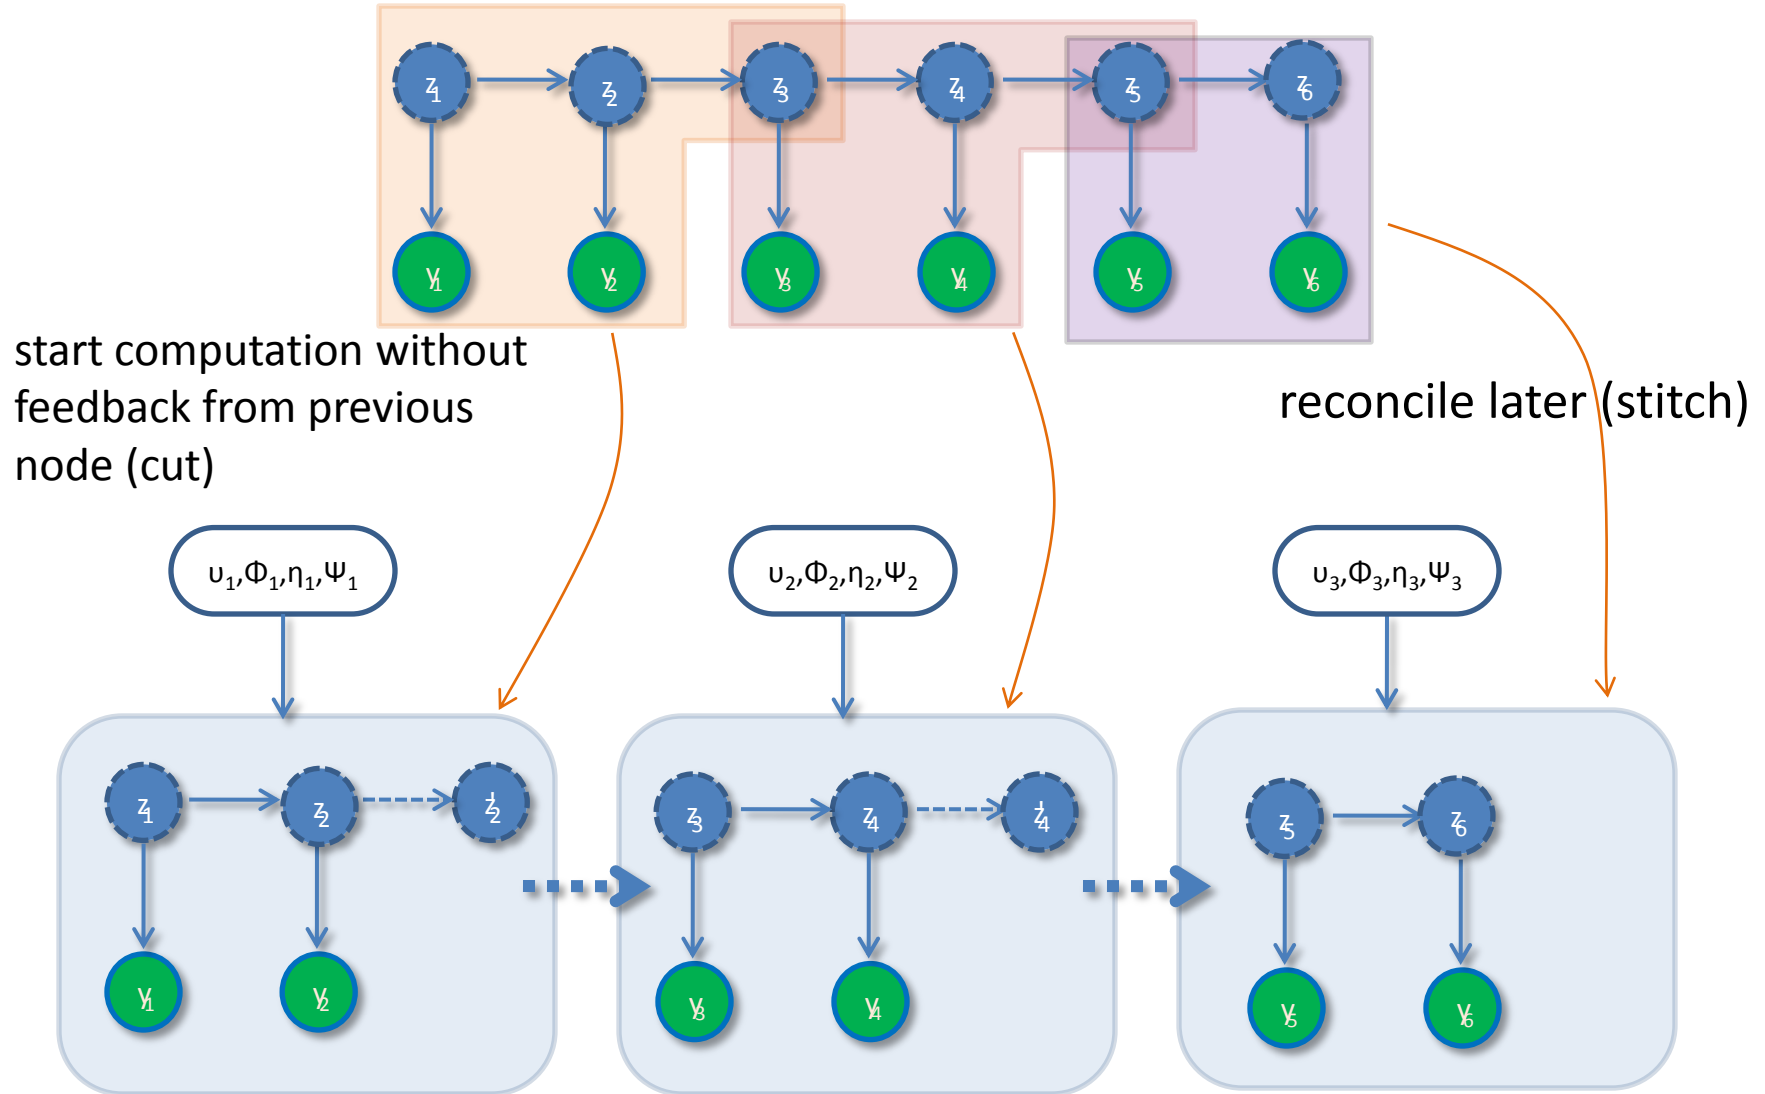


## “Leap of faith”

start computation without feedback from  
previous node (cut),  
and reconcile later (stitch)

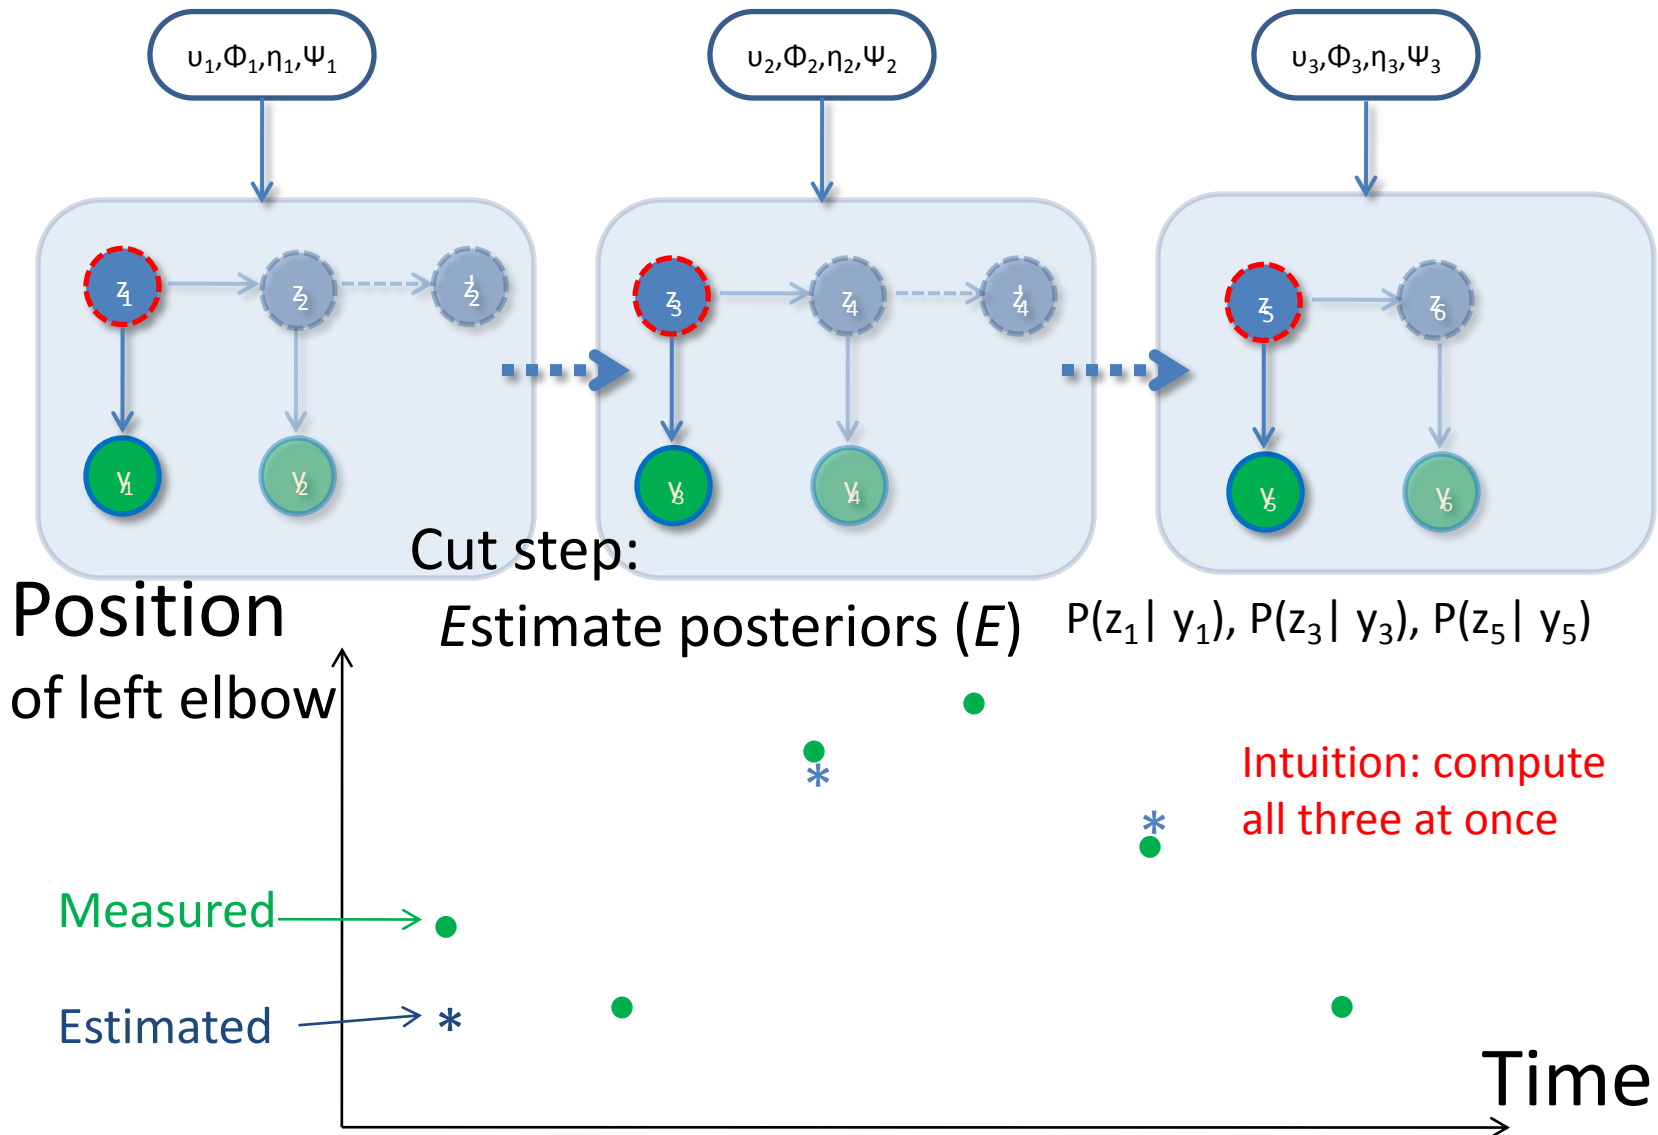


# Proposed Method: Cut-And-Stitch



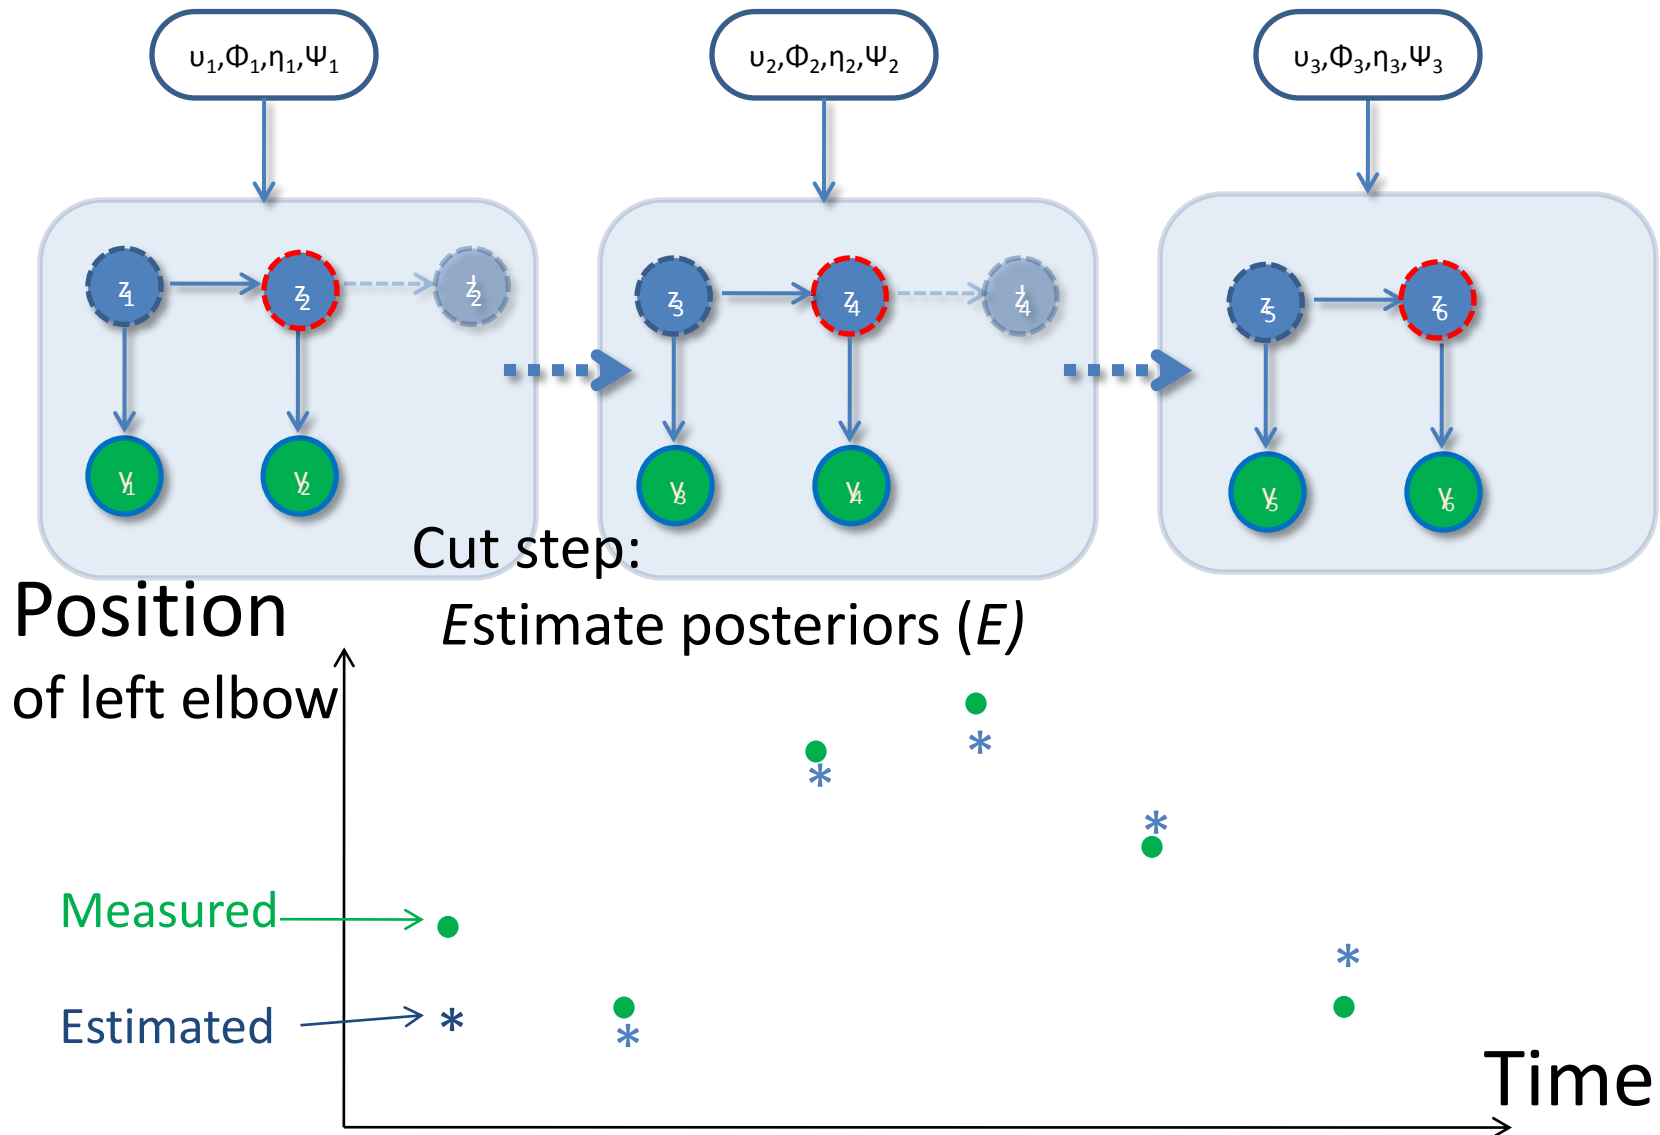


# Cut-And-Stitch



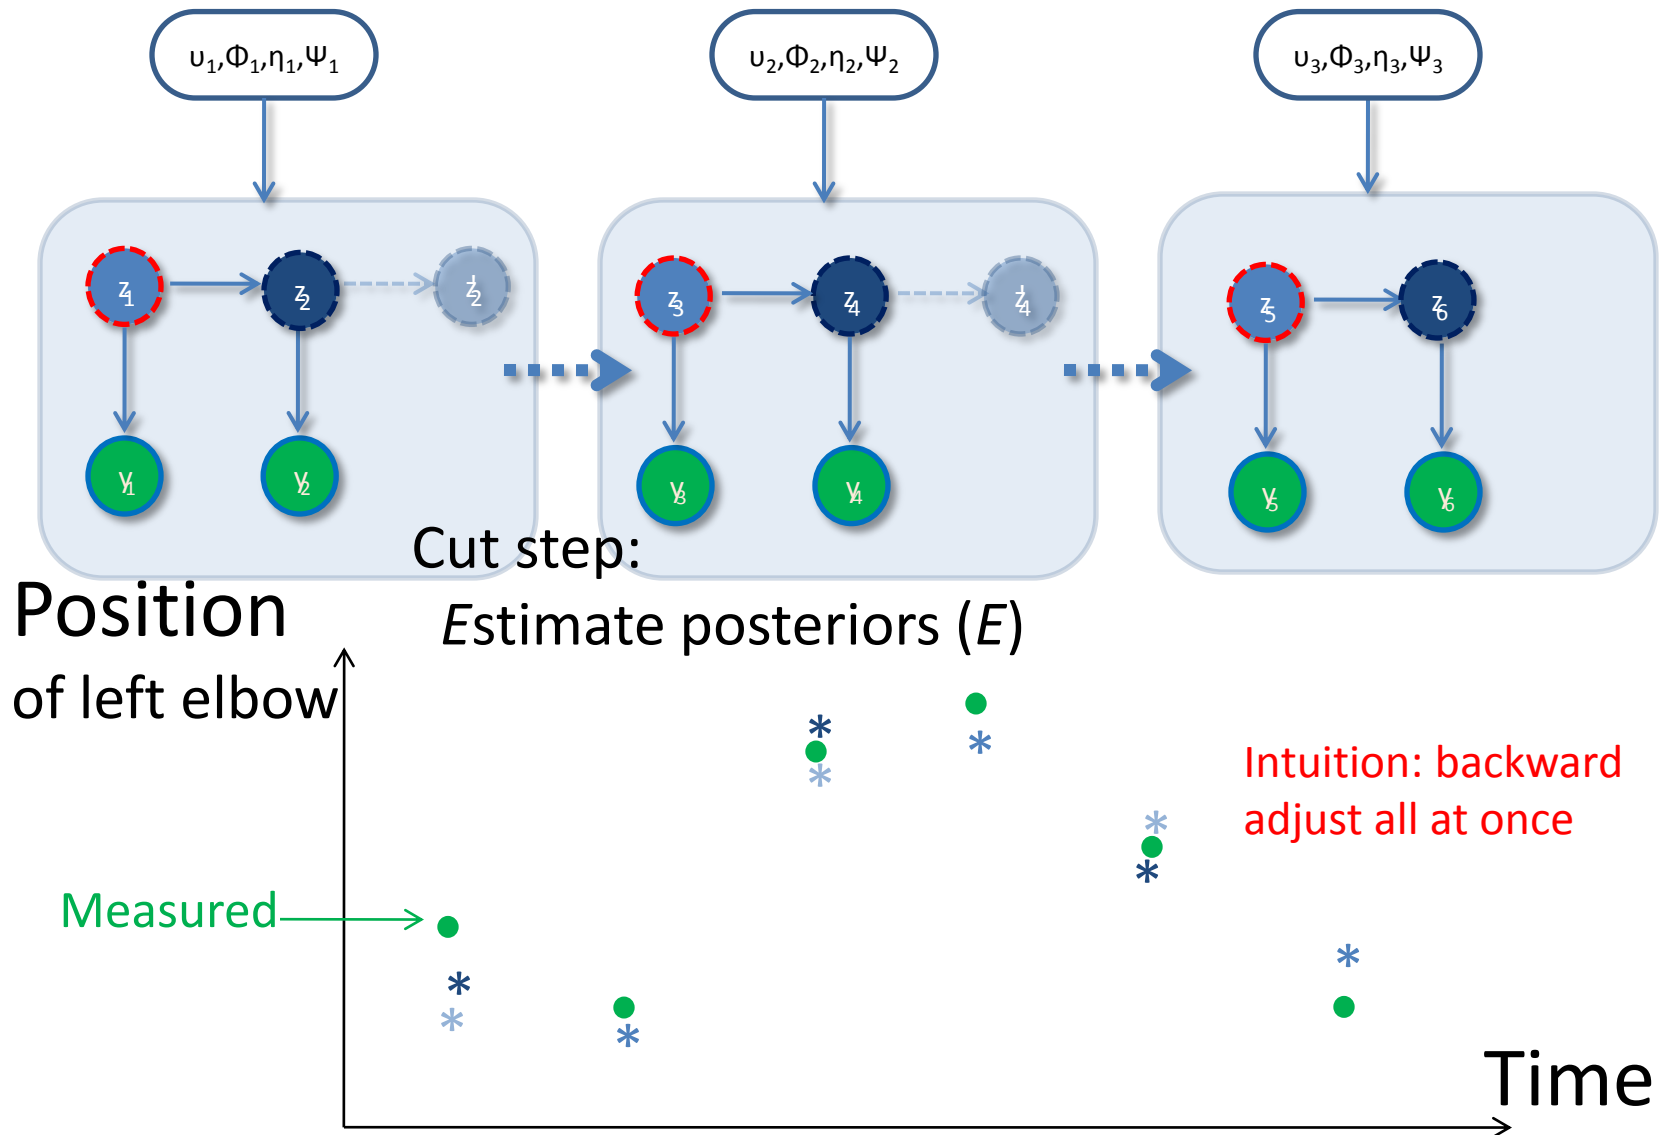


# Cut-And-Stitch



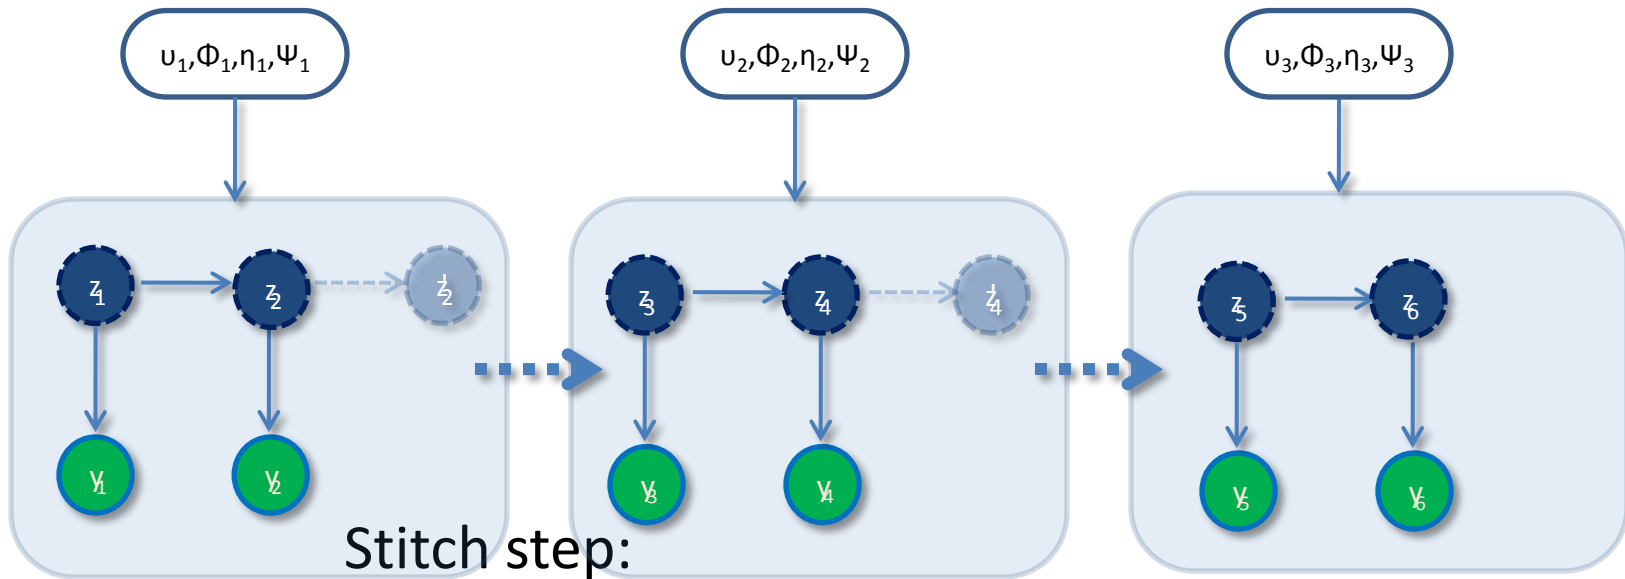


# Cut-And-Stitch






# Cut-And-Stitch



Stitch step:


Position  
of left elbow

Collect sufficient Statistics ( $C$ ) Maximize parameters ( $M$ )





# Cut-And-Stitch



Position  
of left elbow


Re-estimate block parameters ( $R$ )

Intuition: exchange  
messages cross block

Measured

Iterate...  
reconstructed signal

Time







# Outline

- Background
- Motivation: effortless motion stitching
- Parallel learning with Cut-And-Stitch
- Experiments and Results
- Conclusion



# Experiments

Q1: How much speed up can we get?

Q2: How good is the reconstruction accuracy?




# Experiments

- Dataset:
  - 58 human motion sequences, 200 – 500 frames
  - Each frame with 93 bone positions in body local coordinates
  - <http://mocap.cs.cmu.edu>
- Setup:
  - Supercomputer: SGI Altix system, distributed shared memory architecture
  - Multi-core desktop: 4 Intel Xeon cores, shared memory
- Task:
  - Learn the dynamics, hidden variables and reconstruct motion




# Q1: How much speed up?






# Q1: How much speed up?





## Q2: How good?



Result: ~ IDENTICAL accuracy



# Conclusion & Contributions

- A distance function for motion stitching
  - Based on first principle: minimize effort
- General approximate parallel learning algorithm for LDS
  - Near linear speed up
  - Accuracy (NRE):  $\sim$  identical to sequential learning
  - Easily extended to HMM and other chain Markovian models
- Software (C++ w. openMP) and datasets:  
[www.cs.cmu.edu/~leili/paralearn](http://www.cs.cmu.edu/~leili/paralearn)



# Promising Extensions

- Extension
  - HMM
  - other Markov models (*similar graphical model*)
- Open Problem:
  - Can prove the error bound?



# Thank you

- Questions